

# W4

BLOCK 1 · CLASS 4 · PRACTICAL

# Rules, judge, loop.

Last week you watched a feature get built. Today you build one.

2026-09-19

Rules, judge, loop

## W4

BLOCK 1 · CLASS 4 · PRACTICAL

## Rules, judge, loop.

Last week you watched a feature get built. Today you build one.

Timing: open + check-in 6 · what an agent is 9 · three standards 15 · git standards 7 · live kick-off 11 · your turn 30 · wrong/retro/wrap 9 = 87, plus slack. TA-run. Slides carry ideas only: every command is on the Week 4 page, open on their laptops, and in these notes. Presenting laptop: SplitIt with feature 1 merged, clean main, Codex open, GitHub Desktop signed in, data reset.

# Today.

- 1 Check-in
- 2 What an agent actually is
- 3 Three standards
- 4 Git standards
- 5 Build one, live — then you build one
- 6 Retro: one new rule

2026-09-19

Rules, judge, loop

└ Today.

Today.

- 1 Check-in
- 2 What an agent actually is
- 3 Three standards
- 4 Git standards
- 5 Build one, live — then you build one
- 6 Retro: one new rule

# Show me your GitHub page.

- A branch for feature 1.
- A merged pull request.

Not there yet? Start there today.

2026-09-19

Rules, judge, loop

Show me your GitHub page.

Walk the room looking at browsers, not code. Two minutes per row; note who is behind. They can pull main and run the feature-1 tests while you walk — command on the Week 4 page. Don't debug git at the front: point at last week's "when git is confusing" slide.

CHECK-IN · 5 MIN

Show me your GitHub page.

- A branch for feature 1.
- A merged pull request.

Not there yet? Start there today.

## WHAT IS AN AGENT?

# A chat model writes. An agent **acts**.

### A CHAT MODEL

#### **Text in, text out.**

It has never seen your folder.

### AN AGENT

#### **The same model, with hands.**

Reads. Edits. Runs. Sees the result.  
Chooses again.

**Agent = model + tools + a loop.**

2026-09-19

Rules, judge, loop

└ A chat model **writes**. An agent **acts**.

One minute, from their own experience: in Class 2 they asked Codex to show a line of a data file and it showed the real line — a browser chat could not have. Same model family, different powers. "Agentic" just means: can act, not only answer.

WHAT IS AN AGENT?

**A chat model writes. An agent **acts**.**

A CHAT MODEL

**Text in, text out.**

It has never seen your folder.


AN AGENT

**The same model, with hands.**

Reads. Edits. Runs. Sees the result.  
Chooses again.

Agent = model + tools + a loop.


# The model asks. The harness does.



Read · edit · run. That is the whole vocabulary.


Rules, judge, loop

└─ The model asks. The harness does.



Narrate the live kick-off later in exactly these words — “that’s a run, that’s a read, that’s an edit” — so the diagram and the demo are the same thing. The approval prompt is where the human sits inside the loop. The sting that sets up standard 2: an agent can run the tests and still *\*say\** they passed without showing you.

# The loop has a name: **ReAct**.



## THE HARNESS

Runs the tool. Asks you. Feeds the result back. Decides when to stop. Codex is one; so are Claude Code and Cursor.

## WHY YOU CARE

They all run this same loop. Learn it once, and the tools are interchangeable.

Yao, Zhao, Yu, Du, Shafraan, Narasimhan & Cao (2022). *ReAct: Synergizing Reasoning and Acting in Language Models*. ICLR 2023 · [arxiv.org/abs/2210.03629](https://arxiv.org/abs/2210.03629)

2026-09-19

Rules, judge, loop

└ The loop has a name: **ReAct**.

REASON · ACT · OBSERVE

The loop has a name: **ReAct**.

### THE HARNESS

Runs the tool. Asks you. Feeds the result back. Decides when to stop. Codex is one; so are Claude Code and Cursor.


### WHY YOU CARE

They all run this same loop. Learn it once, and the tools are interchangeable.

Yao, Zhao, Yu, Du, Shafraan, Narasimhan & Cao (2022). *ReAct: Synergizing Reasoning and Acting in Language Models*. ICLR 2023 · [arxiv.org/abs/2210.03629](https://arxiv.org/abs/2210.03629)

ReAct = "Reasoning and Acting", Yao et al. 2022 (ICLR 2023, arXiv:2210.03629): a model does far better alternating a thought with an action and reading the result than answering in one go. Every harness since is that loop with better plumbing. The line to land: the three standards are you configuring the harness — the rulebook it loads, the judge it runs, the stop condition.

# Eight years, four kinds of capability.




The job changed twice: when it learned to take **instructions**, and when it got **tools and a loop**.

2026-09-19

Rules, judge, loop

└ Eight years, four kinds of capability.

Two minutes, one sweep left to right. Say the four band names as you go — completing, conversing, acting, persisting — that is the shape to remember, not the model names. If asked for an example of reasoning: an older model asked why the balances were wrong guesses a plausible edit; a reasoning model works the arithmetic out and lands on the actual line. Dates are approximate release months on purpose.



3

THE CORE OF THE CLASS

# Three standards.

A rulebook it reads. A judge it cannot argue with. A loop that knows when to stop.

2026-09-19

Rules, judge, loop

3

THE CORE OF THE CLASS

## Three standards.

A rulebook it reads. A judge it cannot argue with. A loop that knows when to stop.

# AGENTS.md — read before anything else.

## WHAT IT IS

### A README for agents.

Every session starts by reading it.

## WHAT GOES IN

### What it cannot guess.

Commands. Boundaries. How to report back.

## WHAT STAYS OUT

### What it can read.

Your code, explained. Every extra line hides a real one.

It lives in git. When the agent misbehaves, fix a line here — not the prompt.

2026-09-19

Rules, judge, loop

AGENTS.md — read before anything else.

Ask Codex on stage: "What does AGENTS.md tell you to do?" — it recites the file. That is the proof it reads it. The test for every line: would removing it cause a mistake? If not, cut it. Mention the 32 KB cap only if asked; the real cap is attention.

AGENTS.md — read before anything else.

WHAT IT IS  
A README for  
agents.  
Every session starts by  
reading it.

WHAT GOES IN  
What it cannot  
guess.  
Commands. Boundaries.  
How to report back.

WHAT STAYS OUT  
What it can read.  
Your code, explained.  
Every extra line hides a  
real one.

It lives in git. When the agent misbehaves, fix a line here — not the prompt.

# Three rules. One commit.

## BOUNDARY

### One branch per spec.

Never straight to `main`.

## EVIDENCE

### Show me the test output.

Not a sentence saying it passed.

## OFF-LIMITS

### Never touch the tests.

If one looks wrong, say so and stop.

Exact wording on the Week 4 page. Add them, read the diff, commit, push.

2026-09-19

Rules, judge, loop

└─ Three rules. One commit.

Do it on stage: ask Codex to add the three rules under "How to answer" in AGENTS.md, read the diff aloud, accept, commit in GitHub Desktop. The third already existed in spirit — now it is explicit, and the student enforces it. Codex only reads AGENTS.md at the start of a thread, so say "new thread" now.

STANDARD 1 • LIVE • 4 MIN

Three rules. One commit.

BOUNDARY

One branch per spec.

Never straight to `main`.

EVIDENCE

Show me the test output.

Not a sentence saying it passed.

OFF-LIMITS

Never touch the tests.

If one looks wrong, say so and stop.

Exact wording on the Week 4 page. Add them, read the diff, commit, push.

# Decide what **done** means. First.

IT STOPS WHEN IT LOOKS DONE

**So give it a check.**

Failing tests are the to-do list. Green is the finish line.

NEVER MOVE THE GOALPOSTS

**Editing a test is cheating.**

Not finishing.

EVIDENCE, NOT CLAIMS

**Output, not adjectives.**

"All tests pass" is a sentence.

GREEN IS NOT DONE

**Read the diff. Click the app.**

App wrong, tests green? You found a missing test.

2026-09-19

Rules, judge, loop

└─ Decide what **done** means. First.

Run the red tests on stage (command on the Week 4 page) and open the test file next to the spec: one test per checklist line. "Green is not done" is the card that pays off today — feature 2's form field is not tested at all, so only clicking the app catches a missing one.

STANDARD 2 • THE JUDGE

Decide what **done** means. First.

IT STOPS WHEN IT LOOKS DONE

**So give it a check.**

Failing tests are the to-do list. Green is the finish line.

NEVER MOVE THE GOALPOSTS

**Editing a test is cheating.**

Not finishing.

EVIDENCE, NOT CLAIMS

**Output, not adjectives.**

"All tests pass" is a sentence.


GREEN IS NOT DONE

**Read the diff. Click the app.**

App wrong, tests green? You found a missing test.

## STANDARD 3 • THE LOOP

Red → build → run → fix → green.



### ONE TASK PER LOOP

One bug, or one spec.

### LET IT LOOP

It reads its own failures.  
You read the result.


### KNOW WHEN TO STOP

Green — or twice round  
on the same failure. Then  
a new thread.

2026-09-19


Rules, judge, loop

Red → build → run → fix → green.



The three parts are the whole lesson: a rulebook, a judge, a loop. Everything else in agentic coding is a variation. Today the loop runs inside one Codex turn and you watch; the same idea one level up is a script that re-feeds the prompt until the judge passes — Part II, Block 3.

# Every good ask has four parts.



Written out in full on the Week 4 page. Copy it; change two names.

2026-09-19

Rules, judge, loop

└ Every good ask has four parts.


Four parts, in every loop prompt they will ever write. Make them name the four back to you before you paste it. New thread before pasting — the AGENTS.md rules from five minutes ago are only read at the start of a thread.

Every good ask has four parts.



Written out in full on the Week 4 page. Copy it; change two names.

# Meet Ralph.



## WHAT IT IS

The same prompt, handed to a fresh agent again and again, until the judge says green.


## WHY IT NEEDS TODAY

Nobody is watching. The rulebook keeps it in bounds; the judge is all that can stop it.

2026-09-19

Rules, judge, loop

Meet Ralph.



Thirty seconds, as a teaser — you are not demoing this today. It is named after Ralph Wiggum: cheerfully, relentlessly wrong, and it gets there anyway. Geoffrey Huntley's original was literally 'while ;; do cat PROMPT.md | agent; done', and his line for why it works is "deterministically bad in an undeterministic world" — bounded, repeatable failure plus a check that cannot lie. The point to land, and the reason it closes this act: the agent has no memory between runs, so the PROMPT and the repo carry all the state. That is exactly why a rulebook and a judge are not housekeeping — they are the only things between you and an agent confidently going the wrong way for an hour. Part II, with the team project.



HOW A TEAM STAYS SANE

# Three habits.

main always works · one job, one branch · push before you close the laptop.

2026-09-19

Rules, judge, loop




HOW A TEAM STAYS SANE

## Three habits.

main always works · one job, one branch · push before you close the laptop.

Seven minutes for the three git slides. They learned repo/commit/branch/push last week; this is the etiquette on top. Every diagram is the same shape they will see in GitHub Desktop.

`main` always works.




Risky work happens on a branch. `main` only ever receives finished, green work.

2026-09-19


Rules, judge, loop

`main` always works.



Why it matters to them: next week a teammate clones main. If main is broken, they are blocked. This is the habit that makes the team project survivable. Point out the merge dot: the branch comes home, it does not live forever.

# One job, one branch.



Name the branch after the task. Small commits — each one a step you could explain out loud.

2026-09-19

Rules, judge, loop

└ One job, one branch.

GIT STANDARD 2 · SMALL AND NAMED


One job, one branch.



Name the branch after the task. Small commits — each one a step you could explain out loud.

Two branches here, two jobs: a feature and a bug fix. The names say what they are without opening them. Bad commit messages to call out if you see them today: "update", "fix", "changes". Good: "Fix group title showing the id".

# Your laptop and GitHub, in a circle.



Push before you close the laptop. Pull before you start the next thing.

2026-09-19

Rules, judge, loop

└─ Your laptop and GitHub, in a circle.

GIT STANDARD 3 · THE ROUND TRIP

Your laptop and GitHub, in a circle.



Push before you close the laptop. Pull before you start the next thing.

The pull request is where a human reads the diff — today that human is you, reviewing the agent. That is the whole reason the PR step exists in this course. "Pull before you start" prevents next week's most common team failure: two people editing an old copy.

# Watch it work. Then check its work.

- 1 Branch. Named after the spec.
- 2 Red first. See what is missing before anything is written.
- 3 New thread. Paste the ask. Then sit on your hands.
- 4 Watch it loop.
- 5 Green — now don't trust it. Read the diff. Click the app.
- 6 Commit, push, pull request, merge.

2026-09-19

Rules, judge, loop

└─ Watch it work. Then check its work.

Let it loop for real; 2-4 rounds is normal. Narrate it as tool calls. Commands and file names are on the Week 4 page if you want them on screen. The two things it commonly misses on this feature: it forgets to refresh the data files from the seed (tests stay green while the app crashes on load) and it never adds the form field (nothing tests that screen). Both are the "green is not done" lesson — show them, do not fix them silently. If it edits a test: stop, discard, point at the rule, re-prompt.

LIVE · FEATURE 2 · 11 MIN

Watch it work. Then check its work.

- 1 Branch. Named after the spec.
- 2 Red first. See what is missing before anything is written.
- 3 New thread. Paste the ask. Then sit on your hands.
- 4 Watch it loop.
- 5 Green — now don't trust it. Read the diff. Click the app.
- 6 Commit, push, pull request, merge.

# One bug on `main`. One feature on a branch.

## SPLITIT

### Bug 2, then feature 2

Removing a member crashes the page.  
Then expense categories.

## TINY CRM

### Bug 3, then feature 2

"Won this month" is always zero. Then  
conversion by source.

Files, test names and the ask: Week 4 page.

2026-09-19

Rules, judge, loop

└ One bug on `main`. One feature on a branch.

Roam for two things: did they ask for the test output, and did they read the diff before accepting. Never answer a failing test — "what does the last line say?". SplitIt bug 002: the test demands a helper that also removes the member's expenses; the assertion message says so. Tiny CRM bug 003 has a twist: stamping the close date is not enough — the month counter gives up on the whole list when any won lead has no date, and the seed has five; the fix is to skip them, not return zero. Anyone still on feature 1 does that first; feature 2 becomes homework.

YOUR TURN · 30 MIN

One bug on `main`. One feature on a branch.

**SPLITIT**  
**Bug 2, then feature 2**  
Removing a member crashes the page.  
Then expense categories.

**TINY CRM**  
**Bug 3, then feature 2**  
"Won this month" is always zero. Then  
conversion by source.

Files, test names and the ask: Week 4 page.

# Five things you will see.

- 1 It edited a test. Discard. Say the rule. Run again.
- 2 "Done" — tests red. Ask for the output.
- 3 Same failure, three times. Stop. Read it yourself. New thread.
- 4 It fixed three other things. Discard. One task per loop.
- 5 Green, app wrong. Write down what the judge missed.

Every line here is a candidate rule for your `AGENTS.md`.

2026-09-19

Rules, judge, loop

└ Five things you will see.

WHEN THE LOOP GOES WRONG

Five things you will see.

- 1 It edited a test. Discard. Say the rule. Run again.
- 2 "Done" — tests red. Ask for the output.
- 3 Same failure, three times. Stop. Read it yourself. New thread.
- 4 It fixed three other things. Discard. One task per loop.
- 5 Green, app wrong. Write down what the judge missed.

Every line here is a candidate rule for your `AGENTS.md`.

# What did it get wrong? Write the rule.

- The worst thing it did today — one sentence.
- Turn it into a rule you could check.
- Commit it to `AGENTS.md` with the reason.

That is how a rulebook grows: from what went wrong, one line at a time.

2026-09-19

Rules, judge, loop

└─What did it get wrong? Write the rule.

Collect the rules on the board; the good ones become the course's AGENTS.md for Block 2. Expect: "refresh the data after changing the seed files", "don't edit live data by hand", "show me the diff before committing". Reject vague ones ("be careful") — ask "how would you know it followed this?".

RETRO · 7 MIN

What did it get wrong? Write the rule.

- The worst thing it did today — one sentence.
- Turn it into a rule you could check.
- Commit it to `AGENTS.md`, with the reason.

That is how a rulebook grows: from what went wrong, one line at a time.

# Next week you build alone.

- Feature 2: green, merged, pushed.
- One new rule in `AGENTS.md`
- Push before you close the laptop.

Rulebook · judge · loop.

2026-09-19

Rules, judge, loop

└─ Next week you build alone.

Nothing to hand in. Next week's check-in: feature 2 merged, and a commit on AGENTS.md. Part II of the standards — walking away: scripted loops, goals, a second agent reviewing — comes with the team project in Block 3.

BEFORE NEXT WEEK

Next week you build alone.

- Feature 2: green, merged, pushed.
- One new rule in `AGENTS.md`.
- Push before you close the laptop.

Rulebook · judge · loop.